dm22s1

Topic 08 : Regression2

Part 01 : Regression - Extra

Dr Bernard Butler

Department of Computing and Mathematics, WIT.
(bernard.butler@setu.ie)

Autumn Semester, 2022

Outline

- Regression as a means of minimising sum of the squared errors
- Regression assumptions - what they mean, how they can be used for validation and model building
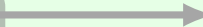- Case studies from Advertising and Credit Balance prediction

Data Mining (Week 8)
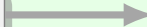
Introduction

Motivating Example

Preparation

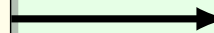Data Handling → Exploring Data 1 → Exploring Data 2 → Building Models
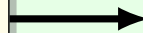
Prediction

Regression 1 → Regression 2 → Classification 1 → Classification 2 → Clustering

Wrap up
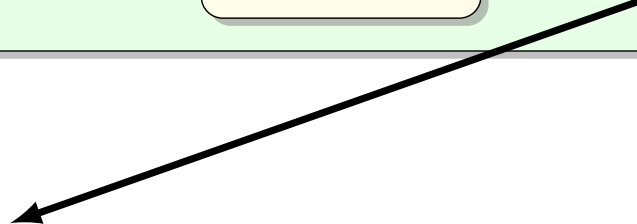
Regression - Extra — Summary

This Week's Aim

This week's aim is to continue the introduction to linear regression, focusing more on how to deal with problems with more challenging datasets.

- Examine some extensions to the simplest case of linear regression.
- We introduce two new concepts: dimensionality reduction and regularisation
- To provide context we will use the following datasets:
  - Generated data (various)
  - Advertising dataset: predicting widgets sold based on spending in different advertising channels
  - Credit dataset: predicting credit balance using income, status, etc.

Case Study 1: Heteroscedasticity - Step 1
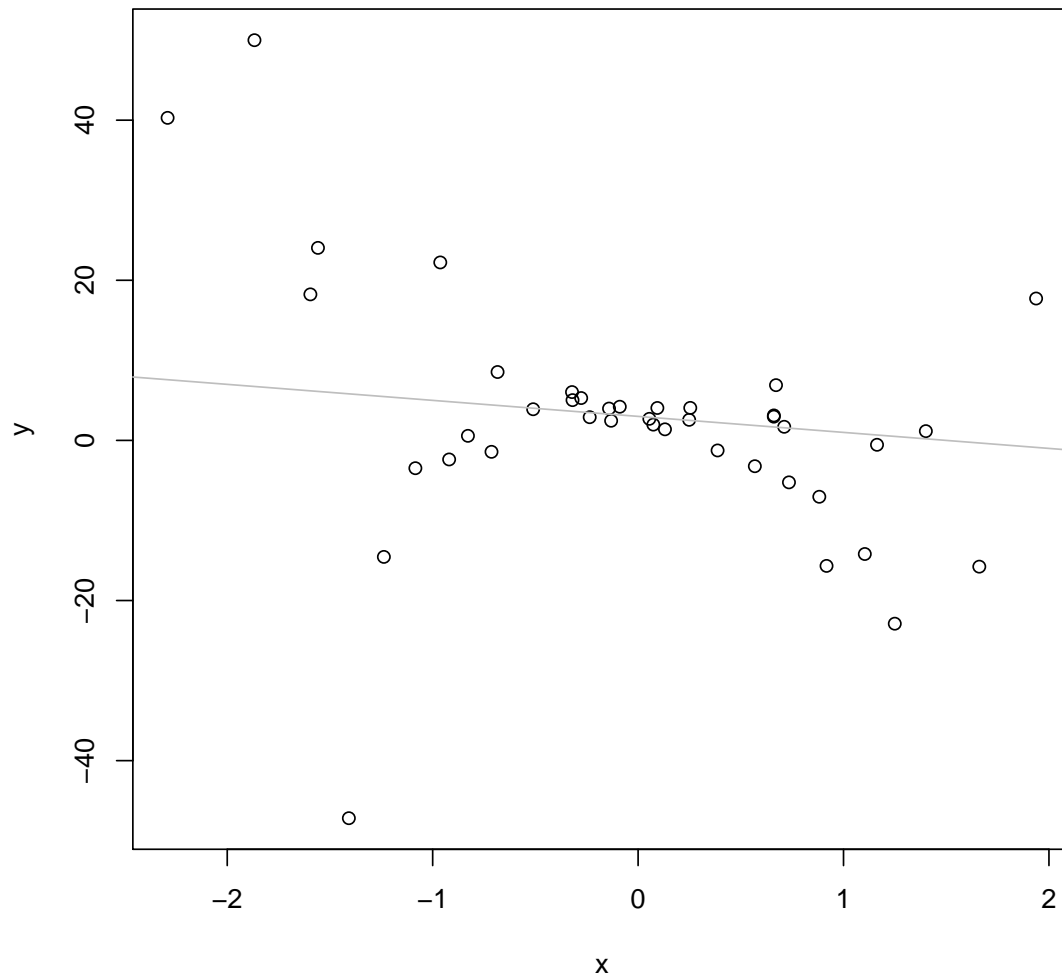
I generated 41 x, y points based on $y = 3 - 2x$, but with added errors that increase away from $x = 0$. The plot shows the line with $\beta = (3, -2)$ in grey.

Case Study 1: Heteroscedasticity - Step 2
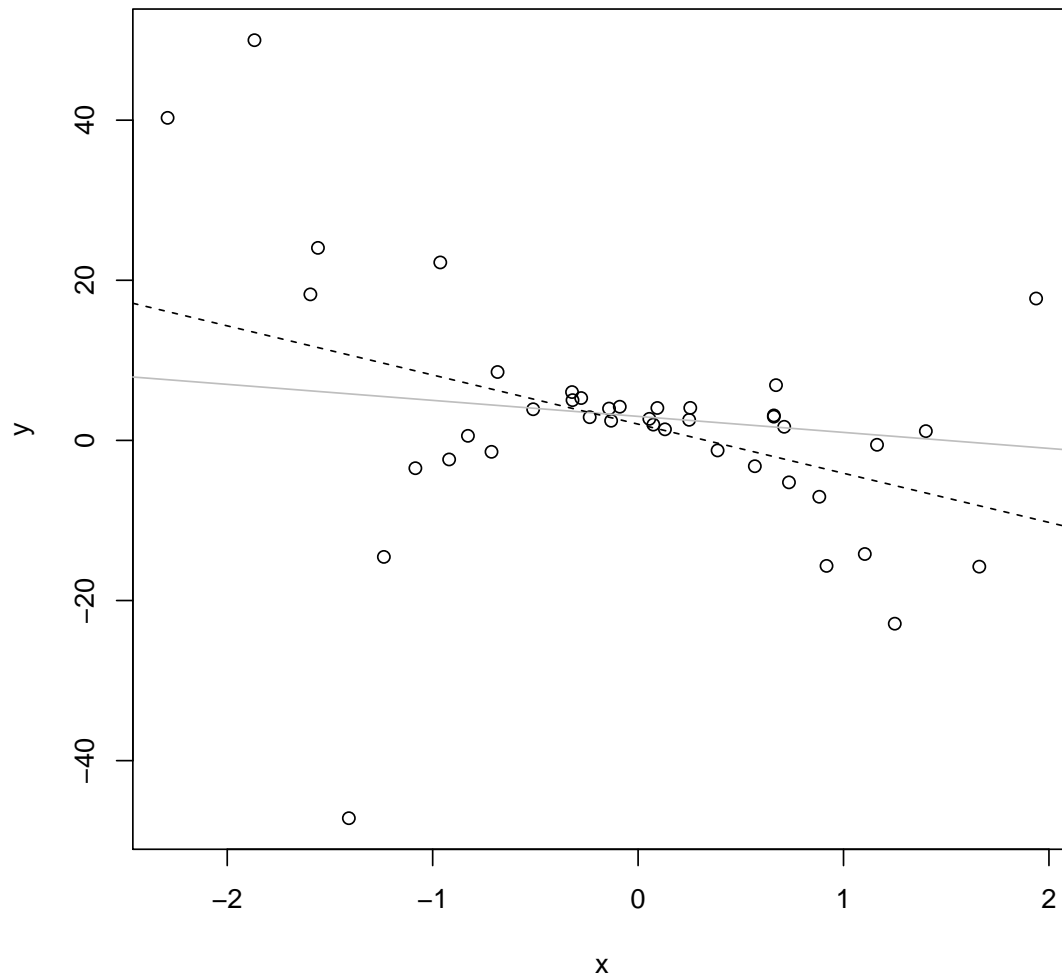
In this plot I added the OLS fit as a dashed line. Note that the parameters of the fit are quite different: $\beta_{OLS} \approx (2, -6)$.

Case Study 1: Heteroscedasticity - Step 3
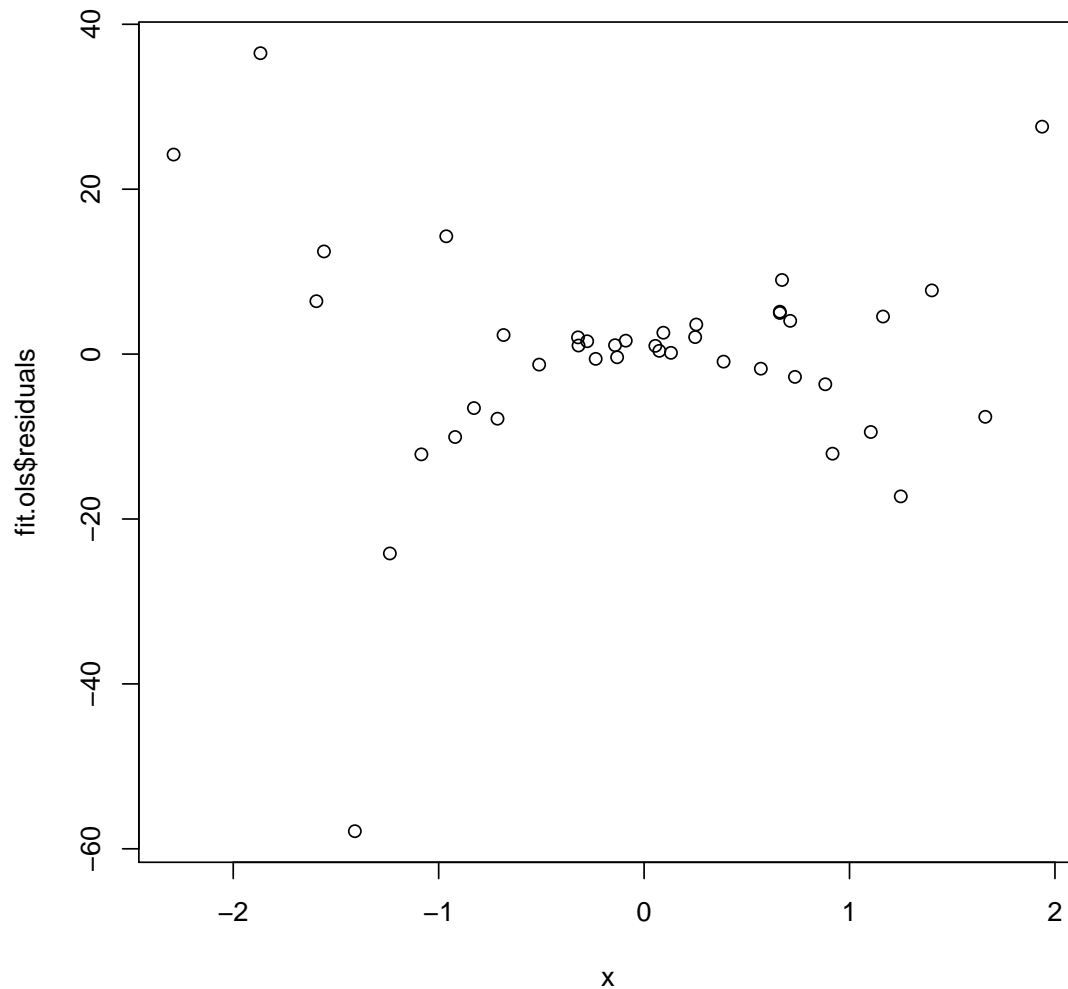
This plot shows how the OLS residuals ϵ_{OLS} increase rapidly away from 0, as expected (since this was how the data was generated).

Case Study 1: Heteroscedasticity - Step 4
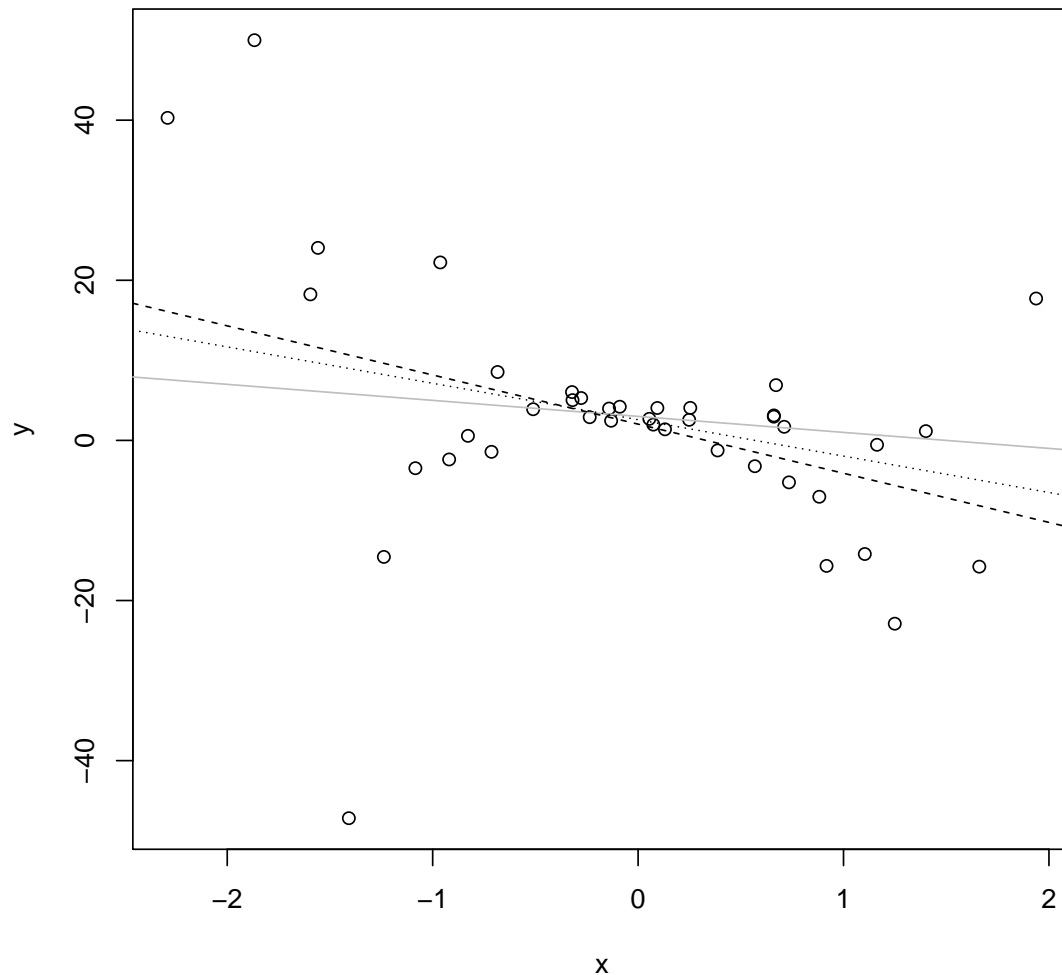
By inspecting the previous residual plot I estimated a weighting function so that the residuals would be "more constant". When this was used to scale the residuals, the resulting Weighted Least Squares estimates were $\beta \approx (2.6, -4.5)$ (shown as a dotted line) and hence closer to the "true" β.

Can you see a problem with finding the weights?

Case Study 3: Advertising: Data and Hypotheses

	TV	Radio	Newspaper	Sales
1	230.1	37.8	69.2	22.1
2	44.5	39.3	45.1	10.4
3	17.2	45.9	69.3	9.3
4	151.5	41.3	58.5	18.5
5	180.8	10.8	58.4	12.9

In this data set, the sales figure captures thousands of widgets of a particular type wss sold in a year. Newspaper, Radio and TV represent the annual spend per widget type on the associated advertising channel. The hypothesis is that spend on advertising is a good predictor of sales performance. Since marketing budgets are limited, where should the adverts be placed for maximum sales?

Case Study 3: Advertising: Looking at the data

Sales vs. Advertising
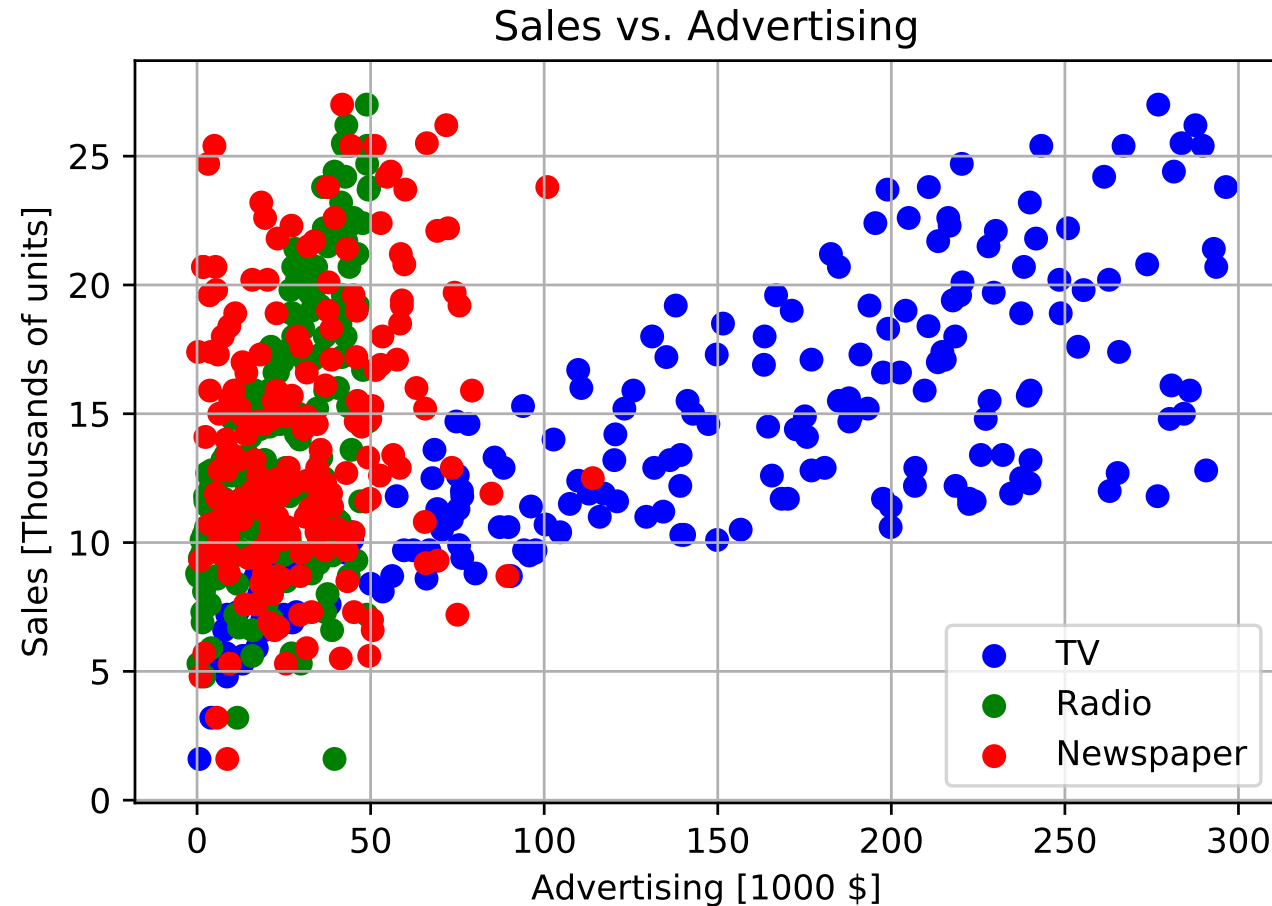
Which of the advertising channels appear to have a linear relationship with Sales?

Case Study 3: Advertising: Collinearity?
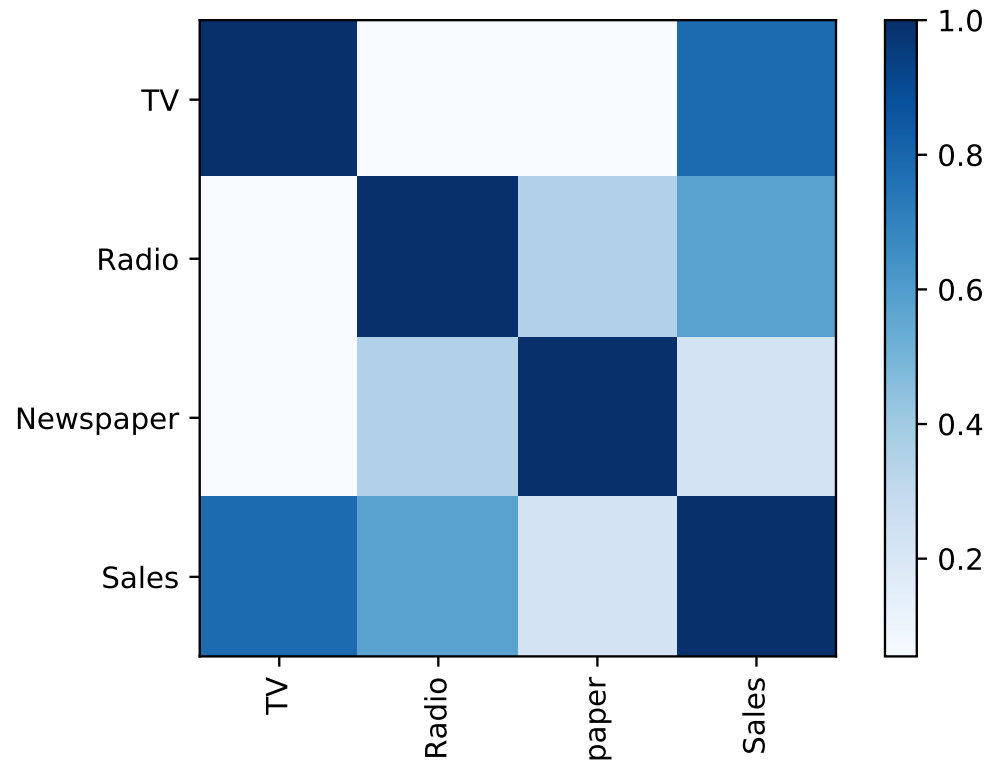
Correlation matrix can indicate which attributes should participate in the model as predictors.

A good predictor should have a high correlation with the dependent variable (Sales in this case) and should have low correlation with other candidate predictors.

What are expected to be good predictors for this data?

Sidebar: specifying models

The statsmodels way

- The dataframe contains the observed variables
- The model is specified separately
- Easier to change the model when experimenting

The sklearn way

- The dataframe contains the (computed) features
- The model is defined implicitly
- Standard interface across all sklearn

statsmodels models look like "Sales $\sim$ TV * Radio + poly(Newspaper,2)"; notation came from applied statistics community.

statsmodels offers its own plotting (like seaborn but not as good). Model summary is very convenient.

sklearn exposes more of the details (e.g., choice of algorithm and configuration parameters).

Both statsmodels and sklearn use the same libraries (scipy, numpy, etc.) underneath.

Case Study 3: Advertising: Model Building

- Start from a "full model" and prune, versus from an "empty model" and add
- We choose the latter, as it is often easier to avoid overfitting

Example 1 (Forward Selection for Advertising Data)

Define: regression model score: (adjusted) R^2 for a given model.

1. Fit "Sales $\sim$ Newspaper", "Sales $\sim$ Radio", "Sales $\sim$ TV" and calculate their R^2 values.

2. Choose the best (largest R^2) single-term model ("Sales $\sim$ TV" in this case) with $R^2 = 0.61$.

3. Fit "Sales $\sim$ TV + Newspaper" and "Sales $\sim$ TV + Radio" and choose the best by R^2 score, which is "Sales $\sim$ TV + Radio" with $R^2 = 0.9$, which is significantly better.

4. Fit "Sales $\sim$ TV + Radio + Newspaper". Its R^2 score is also 0.9, so we favour the existing simpler two-term model (Occam's Razor: other things being equal, choose the simplest model.).

So our preferred model is "Sales $\sim$ TV + Radio" with adjusted $R^2 = 0.9$.

Case Study 3: Advertising: Viewing the Model

Regression: Sales ~ Radio + TV Advertising
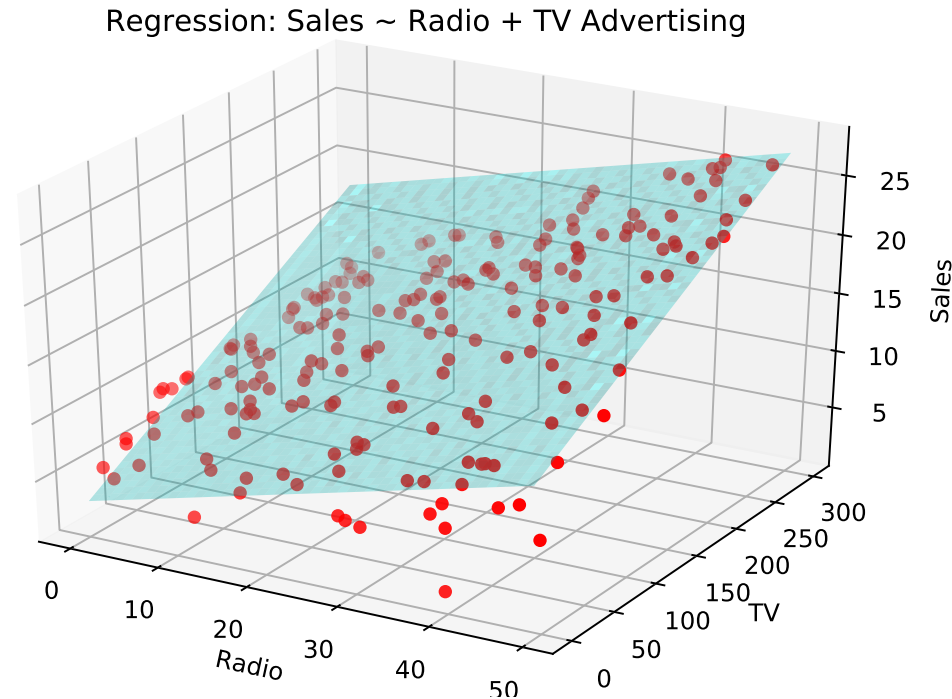
Since this two-term model ignores the contribution of the newspaper channel, the Newspaper spend as a contribution to Sales is just another component of the unmodelled (and apparently random) contribution to Sales.

However, the result is a model where every term is highly significant and the model "explains" 90% of the variance of the data, which is high for an observational study. **Why? Can we do better?**

Case Study 3: Advertising: Interactions; Interpretation

- Trying powers of the Radio and TV greater than 1 did not offer much more.
- However, by adding the TV,Radio interaction so that the model became "Sales $\sim$ TV + Radio + TV:Radio" or equivalently "Sales $\sim$ TV * Radio", R^2 increased to 0.97 from 0.9, which is a significant improvement.
- All β terms have $t-$statistic significance of approximately 0.001 which is extremely significant.
- $\beta_0 = 6.75$, $\beta_{\text{TV}} = 0.019$, $\beta_{\text{Radio}} = 0.029$ and $\beta_{\text{TV:Radio}} = 0.001$, indicating that there is a favourable relationship between TV and Radio advertising ($\beta_{\text{TV:Radio}} > 0$), and that additional spending on Radio results in more Sales than the same spending on TV ($\beta_{\text{Radio}} > \beta_{\text{TV}}$).
- Spending on Newspaper advertising should be discontinued as its contribution to Sales is insignificant (indistinguishable from random noise).

Case Study 4: Credit balances - overview

Introducing

- the `sklearn` approach to regression (we used `statsmodels` with the Diamonds and Advertising data)
- non-numeric explanatory variables like gender and ethnicity
- more advanced regression modelling, e.g., handling correlated variables

Case Study 4: Credit balances - introduction

	Income	Limit	Rating	Cards	Age	Education	Gender	Student	Married	Ethnicity	Balance
1	14.891	3606	283	2	34	11	Male	No	Yes	Caucasian	333
2	106.025	6645	483	3	82	15	Female	Yes	Yes	Asian	903
3	104.593	7075	514	4	71	11	Male	No	No	Asian	580
4	148.924	9504	681	3	36	11	Female	No	No	Asian	964
5	55.882	4897	357	2	68	16	Male	No	Yes	Caucasian	331

Note the presence of some categorical attributes (Gender, Student, Married, Ethnicity). These can participate in linear regression models to predict a numeric response, but must be coded first.
For example, Gender can become an indicator (0,1)-valued variable of the form IsFemale.
Ethnicity has 3 levels and is replaced by 3-1=2 indicator variables.

Case Study 4: Credit balances - Removing Data

- the purpose of the analysis is to predict credit balances.
- Basic exploratory data techniques (histograms) soon indicated that there were 2 cohorts
  1. those who do not use their cards and/or clear their balance each month
  2. those who use their cards and have nonzero balances
- Removing data relating to the first cohort meant that the remaining data looked more cohesive and also made linear regression easier
- Take-away: look for inconsistent subsets in the data, remove them if possible

Case Study 4: Credit balances - Removing correlated attribute

- Correlations between predictors are relatively high, but that between "Limit" and "Rating" is 1
- Generally, customers with a high rating are allowed to have high credit limits
- Conversely, customers will not be allowed high credit limits unless they have a high credit rating
- "Limit" was removed from the data used for analysis
- Take-away: remove correlated attributes, because they increase the standard error (hence variance) and make the solver's job much more difficult

Case Study 4: Credit balances - Contribution of Categorical Variables
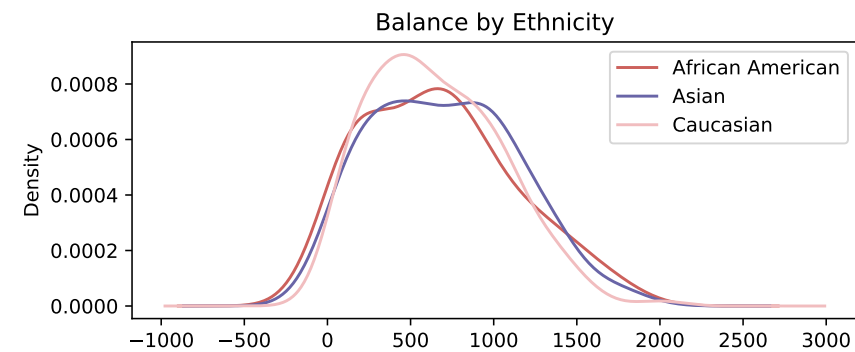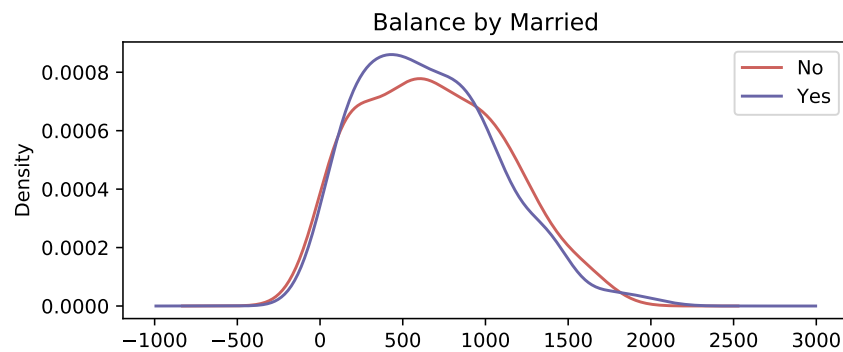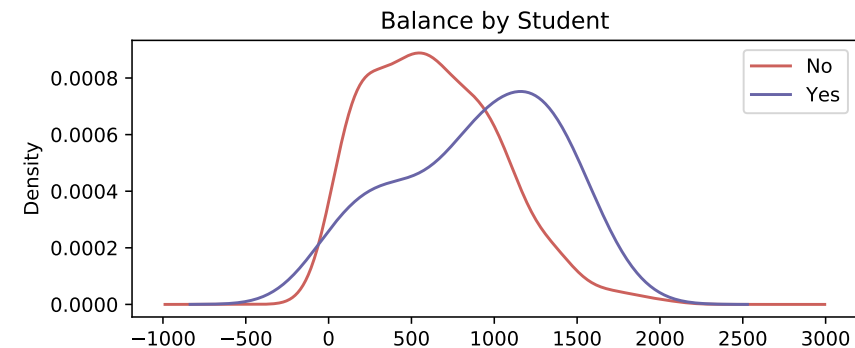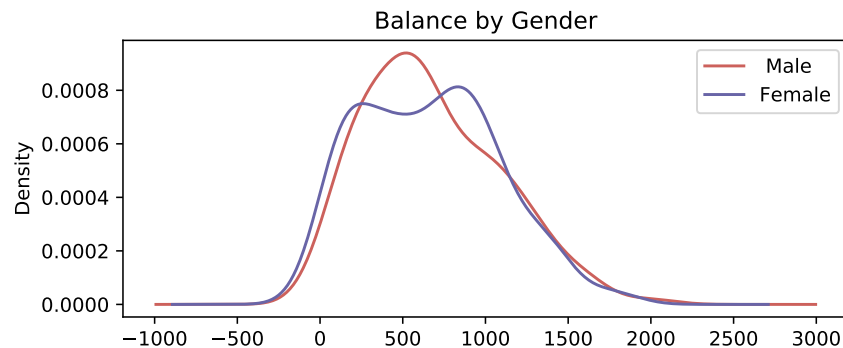
Which of these categorical attributes has a significant effect on Balance?

Case Study 4: Credit balances - Model building

- Using forward selection as before, the best model was found to be "Balance $\sim$ poly(Income,2) + Rating + Age + Student + Income:Rating"

- Could also use Backward Elimination to prune from a complex model

- For this data, high correlations between features can cause difficulties - we need techniques to handle this

Difficulties caused by correlated features

The Problem : Several features are highly correlated, so the solver has difficulty assigning an importance independently to each.

How it shows up : The condition score is large and several model coefficients take large values with opposite signs. Sometimes the solver gives up.

Solution options :

1. Remove selected features from the model (simple, does not always work and requires care)
2. Use *regularisation*, to "penalise" large model coefficients (solve a related problem with a different loss function)
3. Use *dimensionality reduction* (linear PCA) to derive an uncorrelated subset of the features with least loss in explanatory power (principal components can be opaque)

Regularisation introduction

Add *regularisation* constraints to make the model work: $\min_\beta \|\epsilon\|_2^2 + \lambda p(\beta)$

- Options are
    1. *Ridge Regression* where the penalty term takes the form $p(\beta) = \|\beta\|_2$
    2. *Lasso* where the penalty term takes the form $p(\beta) = \|\beta\|_1$
- Regularisation has a metaparameter λ - the challenge is to choose a suitable value
    - if too large: tries less to match the data, increases the bias
    - if too small: tries too hard to match the data so $\beta \to \infty$ and increases the variance

Ridge vs Lasso Regression

Because lasso regression favours the "corners" in parameter space, it tends to set some parameter values to 0 (essentially dropping the associated features). This has the added benefit of making the model smaller and easier to interpret.

Ridge Regression
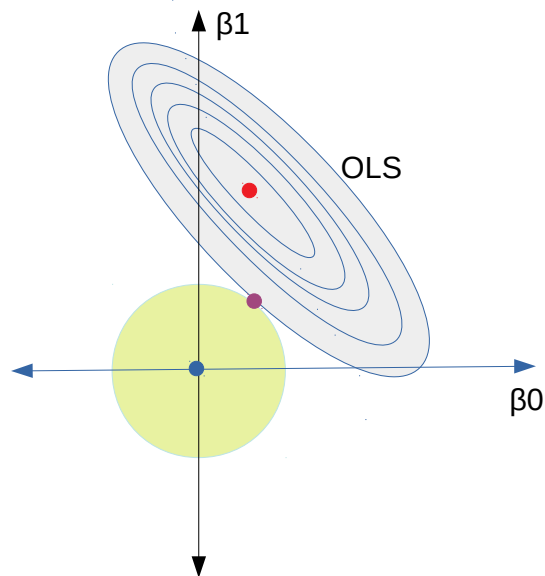
Lasso Regression
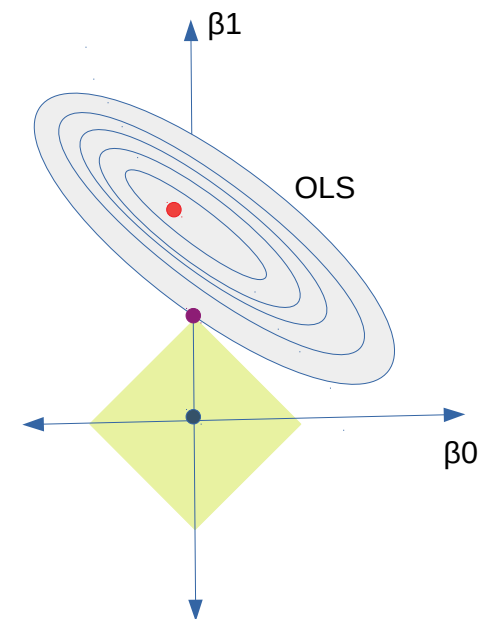
Sidebar - vector norms and their unit balls in 2D
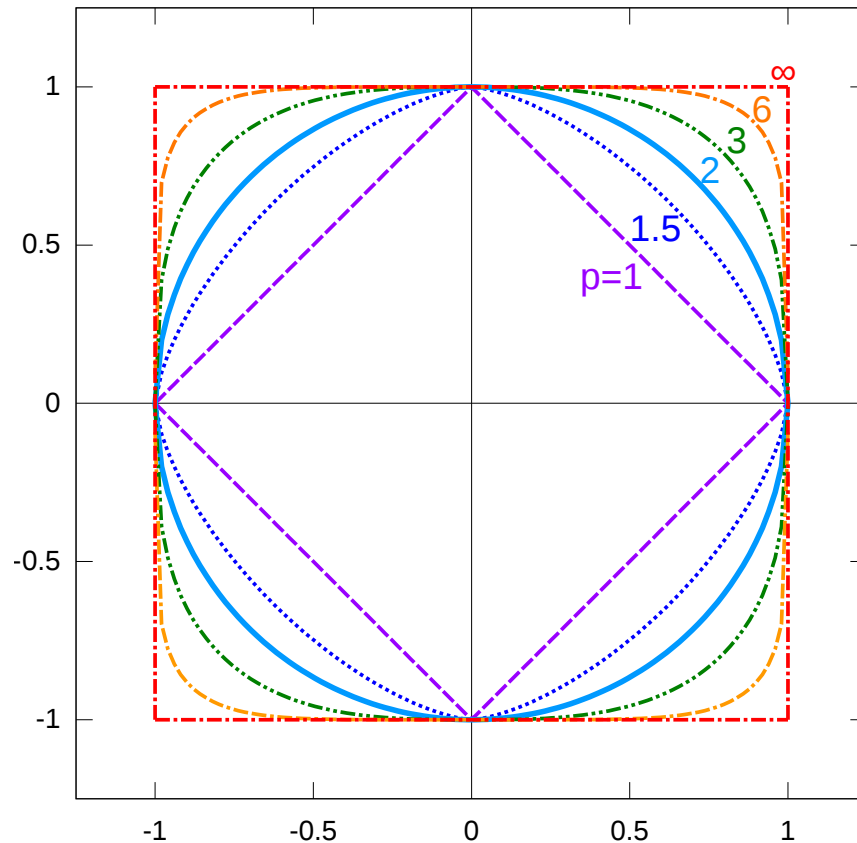
By Quartl - Own work, CC BY-SA 3.0, `https://commons.` `wikimedia.org/w/index.php?curid=17428655`
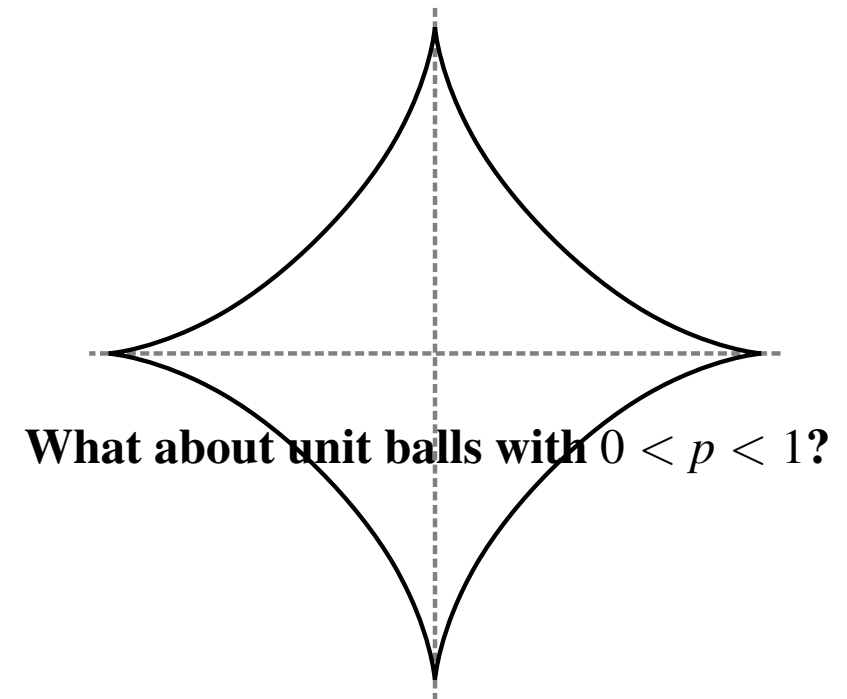
What about unit balls with $0 < p < 1$**?**

Unit ball, $p = \frac{2}{3}$

Public Domain, `https:` `//commons.wikimedia.org/w/` `index.php?curid=1770616`

Case Study 4: Credit balances - Regularisation - Searching for λ

1. Choose a set of candidate λ values

2. For each candidate λ, use K-fold cross-validation on data subsets to estimate the prediction error for the regularised fit with that λ

3. Choose the λ for which the expected error is least

4. Now fit all the training data again with this choice of λ

Note that lasso (but not ridge regression) can set particular β_j to 0 (effectively removing them from the model), so it operates more like backwards elimination in terms of creating a more frugal model having fewer terms.

Ridge regression downweights certain terms but does not set them to zero. However, it can be more performant.

Attribute independence in Multivariate Data

Definition 2 (Covariance)

$\sigma_{12} = E[(X_1 - \mu_1)(X_2 - \mu_2)]$. In words, for two attributes X_1 and X_2, with means μ_1 and μ_2, respectively, σ_{12} is a measure of the the linear dependence between them. If they are independent, we can show that $\sigma_{12} = 0$.

Definition 3 ((Variance-)Covariance Matrix)

When there are n numeric attributes, there are $n \times n$ pairs of covariances $\sigma_{ij}, i = 1, \ldots, n; j = 1, \ldots, n$. The resulting covariance matrix is symmetric and diagonally dominant. This matrix captures the covariance structure of the set of n attributes $\{X_i\}$.

Sometimes it is convenient to work with the correlation matrix, which is a scaled version pf the covariance matrix, with elements $\rho_{ij} = \frac{\sigma_{ij}}{\sigma_i \sigma_j}$, which is scaled so that all the diagonal elements are 1 and the off diagonal elements satisfy $-1 < \rho_{ij} < 1$. If two attributes are highly correlated, adding the second into the model does not increase the explanatory power of the model. Therefore, it pays to determine the covariance matrix from the data before building any models.

Multivariate data with correlated measurements

Example 4 (Measles cases, by city, per week from 1948–1985)

This data spans the period before and after the introduction of vaccination for measles (during the mid 1960s). Measles cases are recorded per week in 7 English cities. Although the cities are not adjacent, it is likely that there will be some spatial autocorrelation. Also, by the nature of disease outbreaks, there will also be some temporal autocorrelation per city.

Row ID	Date	London	Bristol	Liverp...	Manch...	Newc...	Birmin...	Sheffield
Row0	1948–01–10	?	3	40	22	58	78	9
Row1	1948–01–17	240	4	51	19	52	84	11
Row2	1948–01–24	284	3	54	23	34	65	11
Row3	1948–01–31	340	5	54	31	25	106	4
Row4	1948–02–07	511	1	89	66	27	142	7
Row5	1948–02–14	649	3	73	60	47	143	3
Row6	1948–02–21	766	13	169	87	46	191	6
Row7	1948–02–28	932	5	212	61	66	208	9
Row8	1948–03–06	1303	4	283	79	57	290	7
Row9	1948–03–13	1257	15	285	56	82	310	10
Row10	1948–03–20	1716	9	279	85	92	425	5

Removing redundant attribues, based on correlation filters

Row ID	D London	D Bristol	D Liverp...	D Manch...	D Newc...	D Sheffield
London	1	0.474	0.295	0.52	0.52	0.539
Bristol	0.474	1	0.228	0.438	0.374	0.68
Liverpool	0.295	0.228	1	0.432	0.482	0.33
Manchester	0.52	0.438	0.432	1	0.554	0.523
Newcastle	0.52	0.374	0.482	0.554	1	0.536
Birmingham	0.707	0.547	0.365	0.472	0.646	0.691
Sheffield	0.539	0.68	0.33	0.523	0.536	1

The Pearson product-moment correlation matrix was computed and filtered on the critical value $\rho^{(\text{crit})} = 0.7$. Because the correlation between London and Birmingham exceeded 0.7, one of them was removed (Birmingham) so its column was omitted from the filtered correlation matrix above.

Note that the correlation matrix based on rank (e.g., Kendall's τ) is different and the linear dependence between London and Birmingham data was much weaker in that case.

Working with high-dimensional data

Definition 5 (Curse of Dimensionality)

High dimensions do not just require more computing resources and make interpretation more difficult. They also make it more difficult to capture data that samples very high dimensional spaces efficiently. It can be shown that, as the dimension d increases, most of the volume of a hypercube is near the corners, not near the centre, where data might be easiest to collect. This makes estimating parameters much more difficult.

The proof is based on the fact that, as the dimension d tends to infinity, the *ratio* of the volume of the maximum hypersphere inscribed inside the hypercube of the same dimension, tends to 0. Thus there are good reasons to prefer low dimension approximations to high dimensional space.

In 2D: imagine the largest circle fitting inside a square; ratio is $\frac{\pi r^2}{4r^2} = \frac{\pi}{4} \approx 0.79$.

In 3D: imagine the largest sphere fitting inside a cube; ratio is $\frac{(4/3)\pi r^3}{8r^3} = \frac{\pi}{6} \approx 0.52$.
More generally

$$\frac{V_{\text{hypersphere}}}{V_{\text{hypercube}}} = \frac{\pi^{d/2}}{d 2^{d-1} \Gamma(d/2)} \to 0 \text{ when } d \to \infty$$

Feature reduction

- Sometimes it is possible to use intuition to reduce the dimension, by omitting selected attributes.

- Another possibility is to look for groups of correlated attributes (c.f., *mediation*), such as the London and Birmingham measles cases above, and just choose 1 of these.

- More generally, there are techniques that search for a subspace with specified dimension d' of the attributes that captures most of the variance of the full set of attributes having dimension d, where $d' < d$ (often $d' \ll d$).

- The best known of these techniques is *Principal Components Analysis* (PCA).

PCA visualisation

**Mapping to 1 Principal Component**
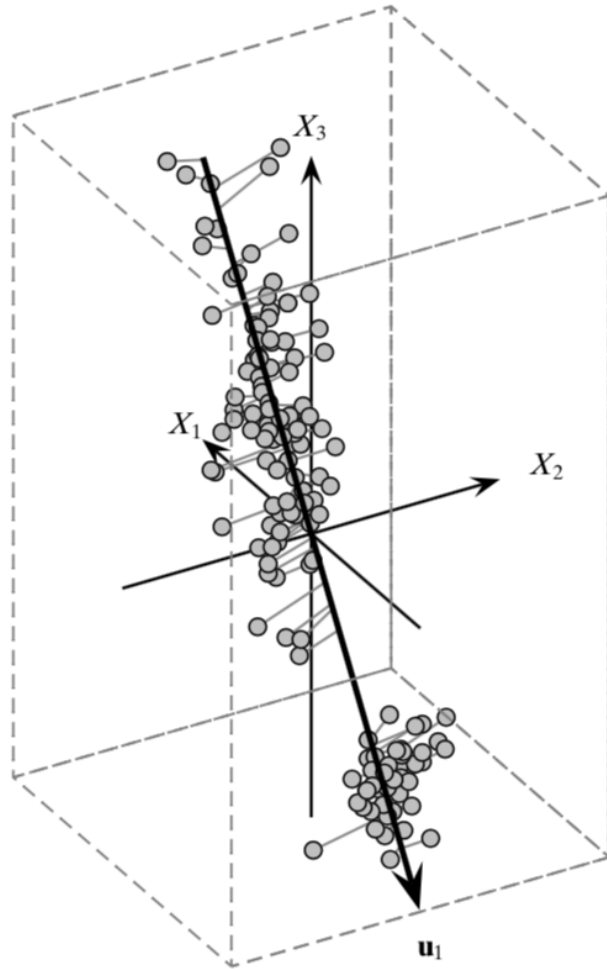
Mapping to 2 Principal Components
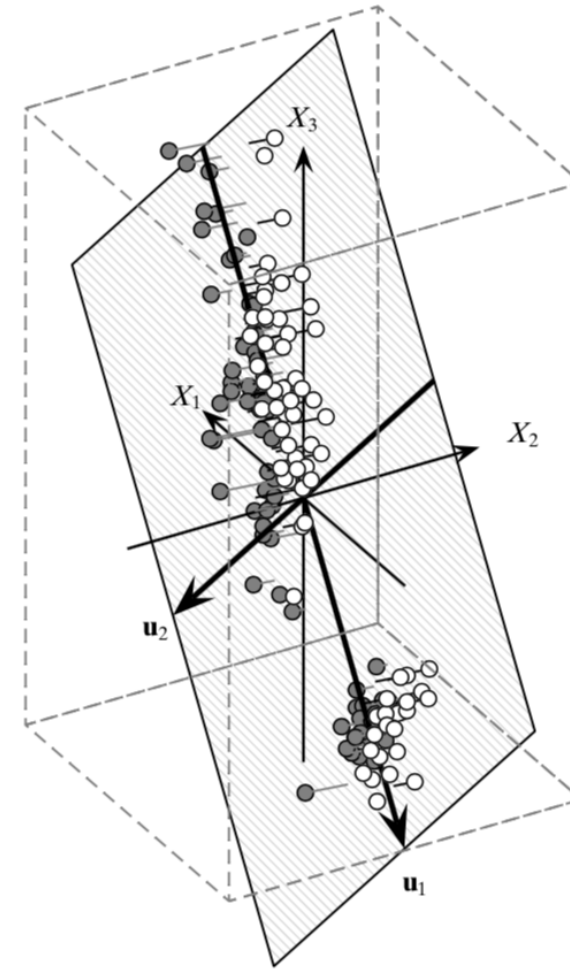
PCA interpretation

- Although the data has dimension $d = 3$, it is possible to find the line (on the left; $d = 1$) and plane (on the right, $d = 2$) which retain most of the variance of the data after it has been projected onto this lower dimensional subspace.

- First compute the transformations needed to align the data with the selected subspace.

- Can then project other data, e.g., test data, onto the subspace that was derived with the original, training data.

- Apply the inverse projection to the data, restoring it to its original orientation. However, because of the use of projections, it is not the same as the original data - the round-trip is "lossy". However, it is helpful to interpret the results in terms of the original attributes.

PCA example
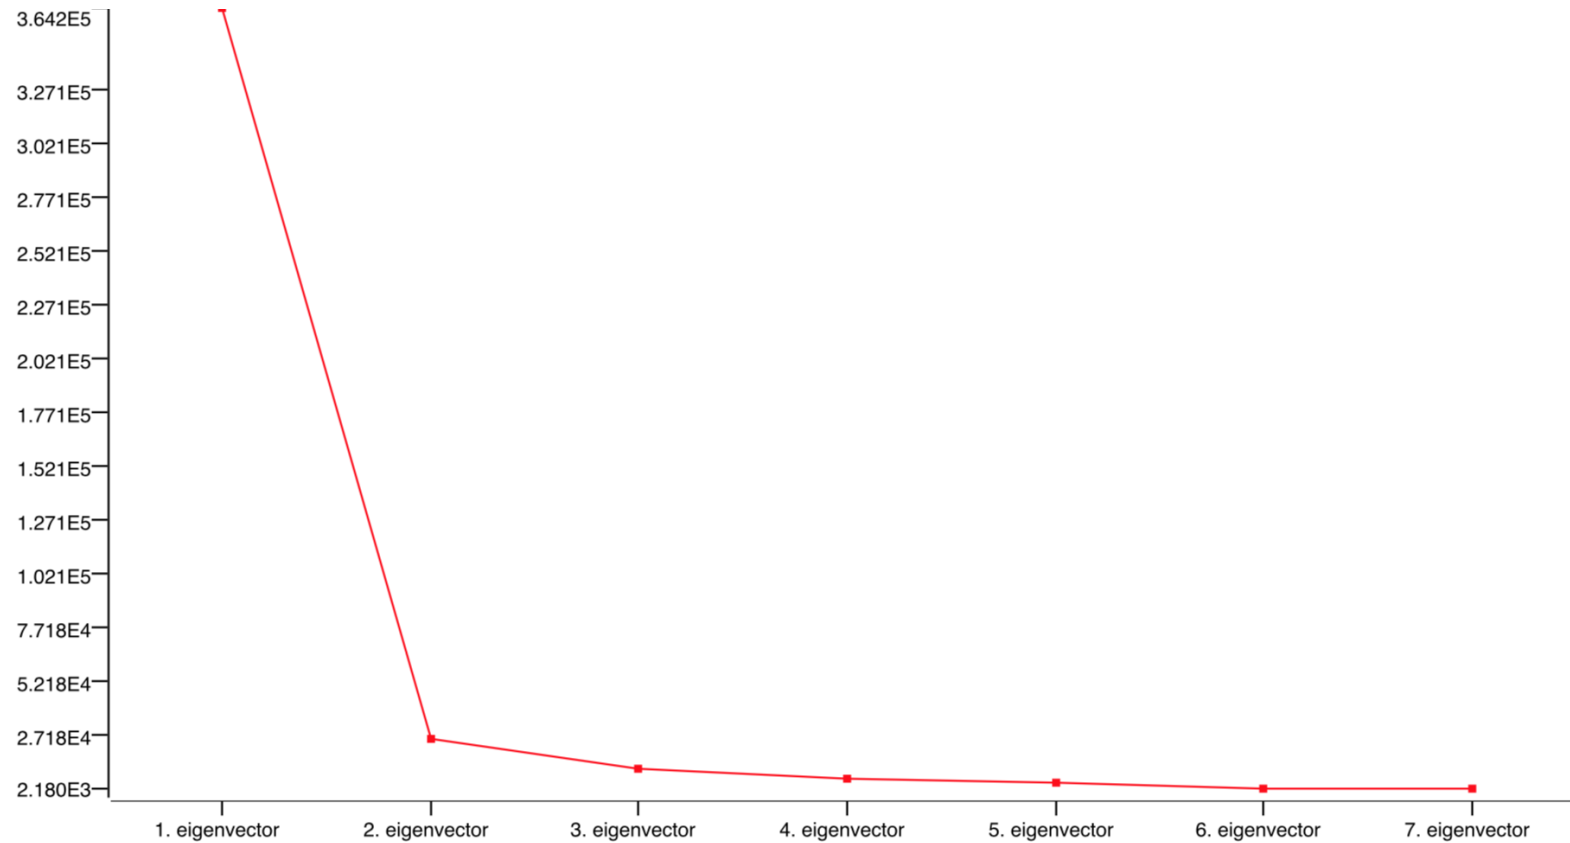
As can be seen from the *scree plot* above, the first eigenvalue captures the bulk of the variance. Arguably, instead of requiring 7 attributes, a single attribute, which is a transformation of the other 7, is sufficient. You could interpret that attribute as representing the measles outbreaks in an archetypal English city…

This youtube video describes PCA concepts well.

Diagnostic plots, statistics and worked examples

See notebooks accompanying this weeks lecture notes.

Review and summary

- Linear regression is one of the foundations of data mining

- It has two phases, of which the first (learning) is generally the most challenging

- It has many variants, so is quite flexible, but flexibility can be abused!

- Careful validation and model building is essential for success - it is an extension of the exploratory work done earlier in the process

- In machine learning, prediction error is the main focus, but you need to be aware of other considerations such as

    1. model parsimony (keep them as small as possible!): faster at both training and evaluation time
    2. the bias-variance dilemma: avoid overfitting and underfitting - remember, your model needs to generalise well from the training to the test set
    3. model interpretability: some models are easier to understand because the terms in the model represent concepts from the domain the data is from

Some Additional Resources

- Book: Introduction to Statistical Learning with R (2013) by James, Gareth and Witten, Daniela and Hastie, Trevor and Tibshirani, Robert.

  I strongly recommend that you read Chapter 3 of the book, as it is very well written and available online for free.

- Kaggle notebooks relating to the datasets addressed this week. There are many, but searching Kaggle should provide nice examples of data mining in action.

- I wrote a report on linear regression that has been added to moodle.